

# Experiment 7:

## Creating a dataset:

```
import pandas as pd  
  
import numpy as np  
  
  
data = {  
  
    "feat1": [5,6,4,5,7,8,5,6,4,5,7,8, 10,11,9,12,13,11,10,12,9,13,12,11,  
              20,21,19,18,22,23,19,20,21,22,18,23, 6,7,5,8,9,7,6,8,5,9,7,8,  
              14,15,13,16,17,15,14,16,13,17,15,16],  
  
  
    "feat2": [10,11,12,9,8,12,11,10,9,8,13,12, 20,22,19,21,23,20,22,21,19,23,22,21,  
              30,29,28,31,32,30,29,31,28,32,30,29, 15,16,14,17,18,16,15,17,14,18,16,17,  
              25,26,24,27,28,26,25,27,24,28,26,27],  
  
  
    "feat3": [2,3,1,2,4,5,2,3,1,2,4,5, 6,7,5,8,9,7,6,8,5,9,7,8,  
              1,2,3,1,4,5,3,2,1,4,3,5, 7,8,6,9,10,8,7,9,6,10,8,9,  
              3,4,2,5,6,4,3,5,2,6,4,5],  
  
  
    "feat4": [1,1,2,1,3,2,1,2,1,3,2,1, 4,4,5,4,6,5,4,5,4,6,5,4,  
              7,8,6,7,9,8,7,8,6,9,8,7, 2,3,1,2,4,3,2,3,1,4,3,2,  
              5,6,4,5,7,6,5,6,4,7,6,5],  
  
}  
  
  
labels = [0]*12 + [1]*12 + [2]*12 + [0]*12 + [1]*12 # 60 labels total  
  
  
df = pd.DataFrame(data)  
df["label"] = labels  
  
  
X = df[["feat1", "feat2", "feat3", "feat4"]]  
y = df["label"]
```

```
df.head()
```

### PCA:

```
from sklearn.decomposition import PCA
```

```
import matplotlib.pyplot as plt
```

```
pca = PCA(n_components=2)
```

```
X_pca = pca.fit_transform(X)
```

```
plt.figure(figsize=(7,5))
```

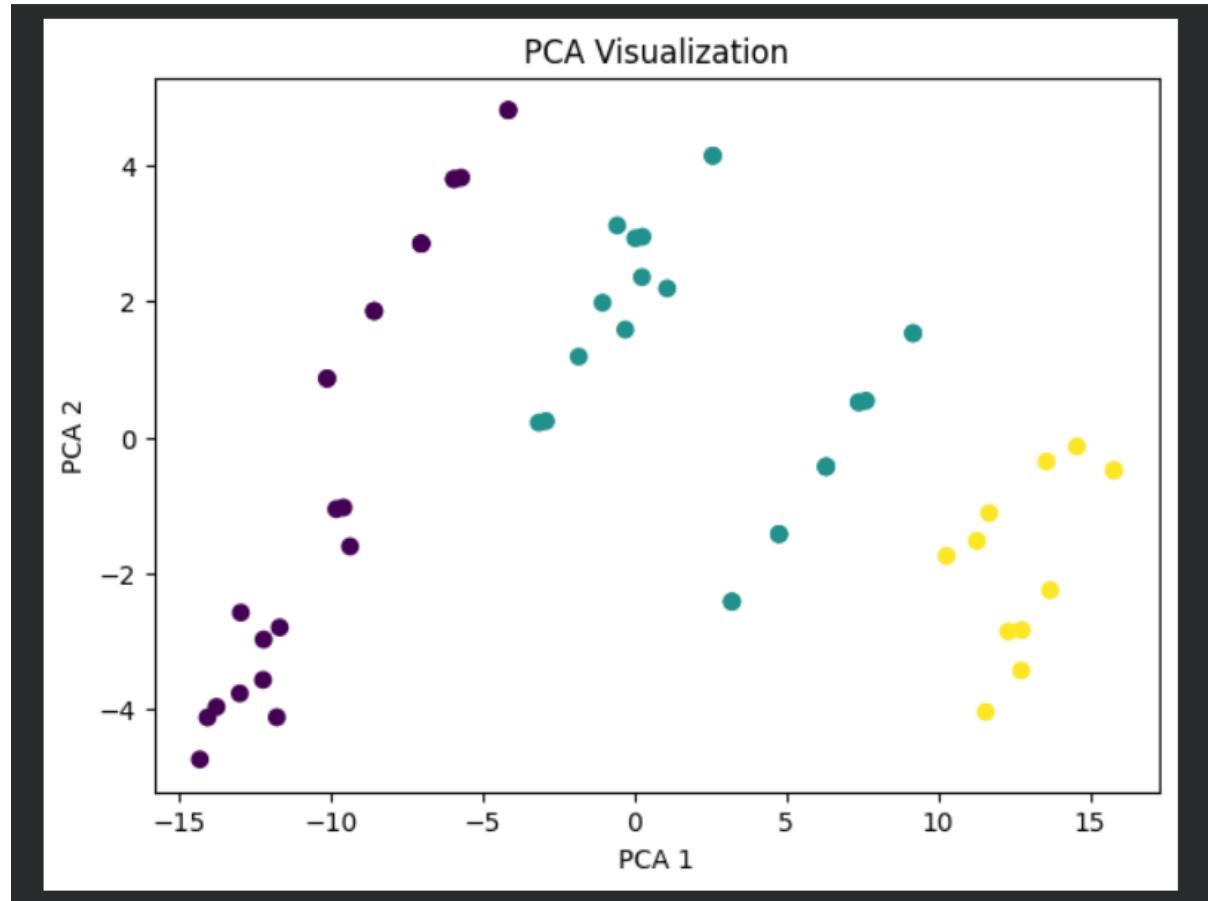
```
plt.scatter(X_pca[:,0], X_pca[:,1], c=y)
```

```
plt.title("PCA Visualization")
```

```
plt.xlabel("PCA 1")
```

```
plt.ylabel("PCA 2")
```

```
plt.show()
```



## LDA:

```
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
```

```
lدا = LinearDiscriminantAnalysis(n_components=2)
```

```
X_lدا = lدا.fit_transform(X, y)
```

```
plt.figure(figsize=(7,5))
```

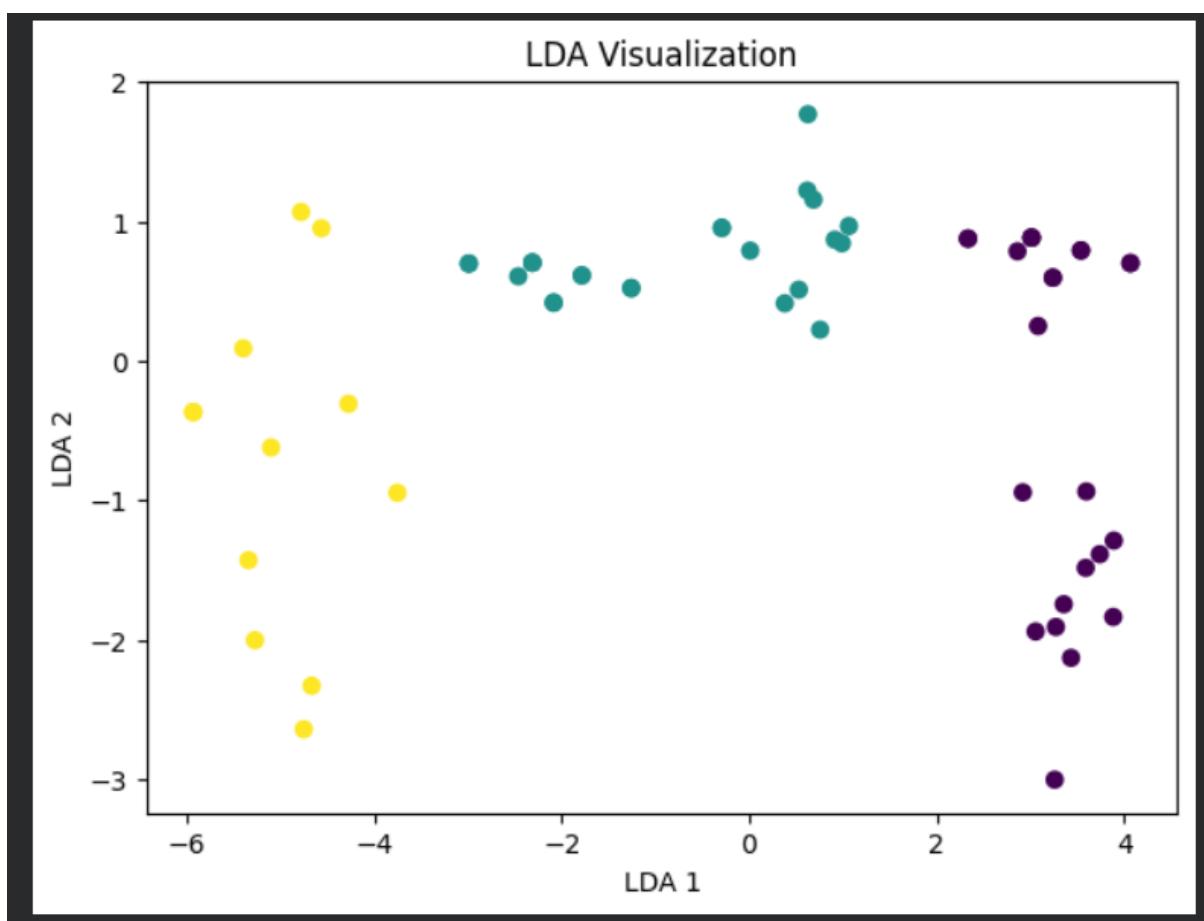
```
plt.scatter(X_lدا[:,0], X_lدا[:,1], c=y)
```

```
plt.title("LDA Visualization")
```

```
plt.xlabel("LDA 1")
```

```
plt.ylabel("LDA 2")
```

```
plt.show()
```



## ICA:

```
from sklearn.decomposition import FastICA
```

```
ica = FastICA(n_components=2, max_iter=500)
```

```
X_ica = ica.fit_transform(X)
```

```
plt.figure(figsize=(7,5))
```

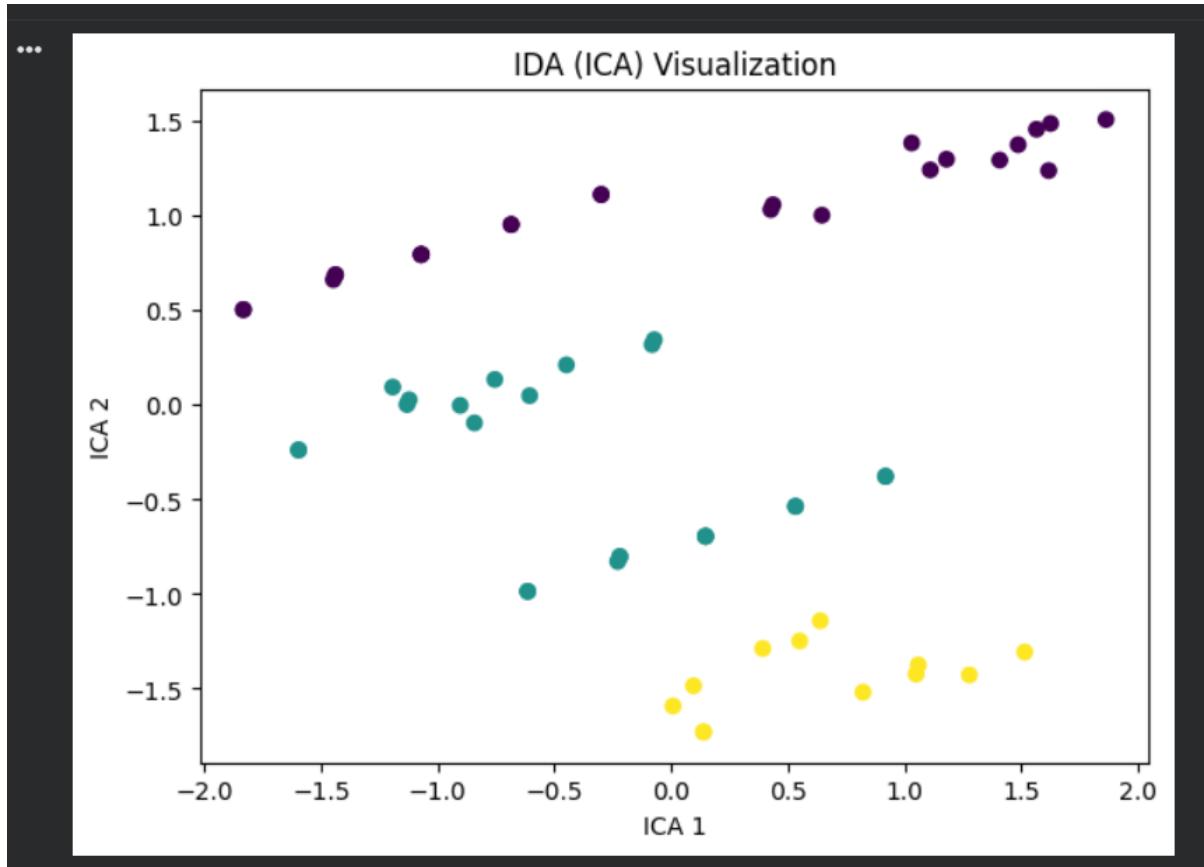
```
plt.scatter(X_ica[:,0], X_ica[:,1], c=y)
```

```
plt.title("IDA (ICA) Visualization")
```

```
plt.xlabel("ICA 1")
```

```
plt.ylabel("ICA 2")
```

```
plt.show()
```



# Experiment 8:

## Simple Linear Regression implementation:

```
import pandas as pd  
  
import numpy as np  
  
import matplotlib.pyplot as plt  
  
from sklearn.linear_model import LinearRegression  
  
from sklearn.model_selection import train_test_split  
  
  
data = {  
    'Experience': [1.1, 1.3, 1.5, 2.0, 2.2, 2.9, 3.0, 3.2, 3.3, 3.7,  
        3.9, 4.0, 4.5, 4.9, 5.1, 5.3, 5.9, 6.0, 6.8, 7.1,  
        7.9, 8.2, 8.7, 9.0, 9.5, 9.6, 10.3, 10.5, 11.0, 11.5],  
    'Salary': [39343, 46205, 37731, 43525, 39891, 56642, 60150, 54445,  
        64445, 57189, 63218, 55794, 56957, 57081, 61111, 67938,  
        66029, 83088, 81363, 93940, 91738, 98273, 101302, 103150,  
        112636, 122391, 121872, 127345, 128765, 135675]  
}  
  
}
```

```
df = pd.DataFrame(data)
```

```
X = df[['Experience']]  
y = df['Salary']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
model = LinearRegression()  
model.fit(X_train, y_train)
```

```
y_pred = model.predict(X_test)
```

```
plt.scatter(X_train, y_train, color='blue', label='Training data')
plt.plot(X_train, model.predict(X_train), color='red', linewidth=2, label='Best Fit Line')
plt.title('Salary vs Experience (Training set)')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.legend()
plt.show()

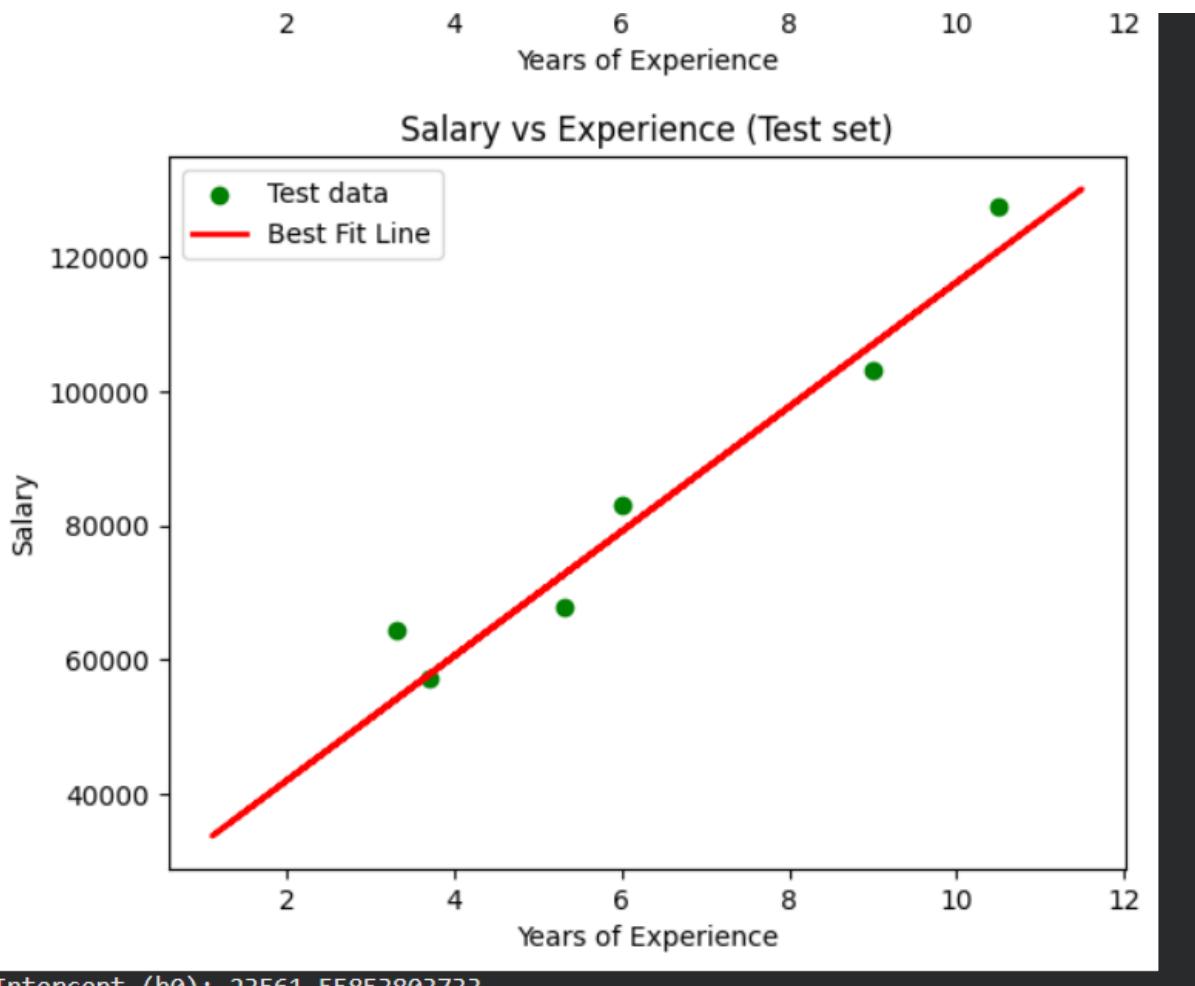
plt.scatter(X_test, y_test, color='green', label='Test data')
plt.plot(X_train, model.predict(X_train), color='red', linewidth=2, label='Best Fit Line')
plt.title('Salary vs Experience (Test set)')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.legend()
plt.show()

print("Intercept (b0):", model.intercept_)
print("Coefficient (b1):", model.coef_[0])
print("\nEquation of line: Salary = {:.2f} + {:.2f} * Experience".format(model.intercept_,
model.coef_[0]))

exp = 5
predicted_salary = model.predict([[exp]])
```

```
print(f"\nPredicted Salary for {exp} years of experience: ₹{predicted_salary[0]:.2f}")
```





## Multiple Linear regression implementation:

```
import pandas as pd

import numpy as np

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.preprocessing import OneHotEncoder

from sklearn.compose import ColumnTransformer

data = {

    'R&D Spend': [165349.2, 162597.7, 153441.51, 144372.41, 142107.34, 131876.9, 134615.46, 130298.13, 120542.52, 123334.88, 101913.08, 100671.96, 93863.75, 91992.39, 119943.24, 114523.61, 78013.11, 94657.16, 91749.16, 86419.7, 76253.86, 78389.47, 73994.56, 67532.53, 77044.01, 64664.71, 75328.87, 72107.6, 66051.52, 65605.48, 61994.48, 61136.38, 63408.86, 55493.95, 46426.07, 46014.02, 28663.76, 44069.95, 20229.59, 38558.51, 28754.33, 27892.92, 23640.93, 15505.73, 22177.74, 1000.23, 1315.46, 0.0, 542.05, 0.0], 'Administration': [136897.8, 151377.59, 101145.55, 118671.85, 91391.77, 99814.71, 147198.87, 145530.06, 148718.95, 108679.17, 110594.11, 91790.61, 127320.38, 135495.07, 156547.42, 122616.84, 121597.55, 145077.58, 114175.79, 153514.11, 113867.3, 153773.43, 122782.75, 105751.03, 99281.34, 139553.16, 144135.98, 127864.55, 182645.56, 153032.06, 115641.28, 152701.92, 129219.61, 103057.49, 157693.92, 85047.44, 127056.21, 51283.14, 65947.93, 82982.09, 118546.05, 84710.77, 96189.63, 127382.3, 154806.14, 124153.04, 115816.21, 135426.92, 154699.78, 127382.3], 'Marketing Spend': [471784.1, 443898.53, 407934.54, 383199.62, 366168.42, 362861.36, 127716.82, 323876.68, 311613.29, 304981.62, 229160.95, 249744.55, 249839.44, 252664.93, 256512.92, 261776.23, 264346.06, 282574.31, 294919.57, 0.0, 298664.47, 299737.29, 303319.26, 304768.73, 140574.81, 137962.62, 134050.07, 353183.81, 118148.2, 107138.38, 91131.24, 88218.23, 46085.25, 214634.81, 210797.67, 205517.64, 201126.82, 197029.42, 185265.1, 174999.3]}
```

```
    172795.67, 164470.71, 157293.19, 155622.91, 152701.92, 140964.55, 134361.69,
130298.13, 123334.88, 120542.52],  
    'State': ['New York', 'California', 'Florida', 'New York', 'Florida', 'California', 'Florida', 'New York',
'California', 'Florida',  
    'New York', 'California', 'Florida', 'New York', 'California', 'Florida', 'New York', 'California',
'Florida', 'New York',  
    'California', 'Florida', 'New York', 'California', 'Florida', 'New York', 'California', 'Florida', 'New
York', 'California',  
    'Florida', 'New York', 'California', 'Florida', 'New York', 'California', 'Florida', 'New York',
'California', 'Florida',  
    'New York', 'California', 'Florida', 'New York', 'California', 'Florida', 'New York', 'California',
'Florida', 'New York'],  
    'Profit': [192261.83, 191792.06, 191050.39, 182901.99, 166187.94, 156991.12, 156122.51,
155752.6, 152211.77, 149759.96,  
    146121.95, 144259.4, 141585.52, 134307.35, 132602.65, 129917.04, 126992.93, 125370.37,
124266.9, 122776.86,  
    118474.03, 111313.02, 110352.25, 108733.99, 108552.04, 107404.34, 105733.54,
105008.31, 103282.38, 101004.64,  
    99937.59, 97483.56, 97427.84, 96778.92, 96712.8, 96479.51, 90708.19, 89949.14,
81229.06, 81005.76,  
    78239.91, 77798.83, 71498.49, 69758.98, 65200.33, 64926.08, 49490.75, 42559.73,
35673.41, 14681.4]  
}
```

```
df = pd.DataFrame(data)
```

```
X = df.iloc[:, :-1]
```

```
y = df.iloc[:, -1]
```

```
ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(drop='first'), [3])],
remainder='passthrough')
```

```
X = np.array(ct.fit_transform(X))
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```

model = LinearRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

results = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
print(results)

print("\nCompany with Maximum Profit:", df.loc[df['Profit'].idxmax(), 'State'])

print("Maximum Profit:", df['Profit'].max())

coef_df = pd.DataFrame({'Feature': ['State1','State2','R&D Spend','Administration','Marketing Spend'], 'Coefficient': model.coef_})

print("\nFeature Impact on Profit:\n", coef_df)

```

	Actual	Predicted
13	134307.35	127535.966981
39	81005.76	83884.741803
30	99937.59	97647.738270
45	64926.08	47893.660005
17	125370.37	128510.806228
48	35673.41	43976.191889
26	105733.54	108662.481700
25	107404.34	101511.446951
32	97427.84	97957.480067
19	122776.86	114300.529646

**Company with Maximum Profit: New York**  
**Maximum Profit: 192261.83**

**Feature Impact on Profit:**

	Feature	Coefficient
0	State1	-1725.065933
1	State2	1069.077047
2	R&D Spend	0.823835
3	Administration	-0.100269
4	Marketing Spend	0.027062

# Experiment 9:

**Simple decision tree:**

**Importing the dataset:**

```
import pandas as pd
```

```
data_set = pd.read_csv('/content/drive/MyDrive/Datasets/user_data.csv')  
print(data_set)
```

Decision Tree algorithm for data related to salary:

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import pandas as pd
```

```
data_set = pd.read_csv('/content/drive/MyDrive/Datasets/user_data.csv')
```

```
x = data_set.iloc[:, [3, 4]].values
```

```
y = data_set.iloc[:, 5].values
```

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=0)
```

```
from sklearn.preprocessing import StandardScaler
```

```
sc = StandardScaler()
```

```
x_train = sc.fit_transform(x_train)
```

```
x_test = sc.transform(x_test)
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
classifier = DecisionTreeClassifier(criterion='entropy', random_state=0)
```

```
classifier.fit(x_train, y_train)
```

```
Prediction and confusion matrix:  
y_pred = classifier.predict(x_test)  
  
print("Predicted:\n", y_pred)  
  
print("\nActual:\n", y_test)  
  
from sklearn.metrics import confusion_matrix  
  
cm = confusion_matrix(y_test, y_pred)  
  
print("\nConfusion Matrix:\n", cm)
```

## Training Set visualization:

```
from matplotlib.colors import ListedColormap  
  
x_set, y_set = x_train, y_train  
  
x1, x2 = np.meshgrid(  
    np.arange(x_set[:, 0].min() - 1, x_set[:, 0].max() + 1, 0.01),  
    np.arange(x_set[:, 1].min() - 1, x_set[:, 1].max() + 1, 0.01)  
)  
  
plt.contourf(  
    x1, x2,  
    classifier.predict(np.array([x1.ravel(), x2.ravel()]).T).reshape(x1.shape),  
    alpha=0.75,  
    cmap=ListedColormap(['purple', 'green']))  
  
plt.xlim(x1.min(), x1.max())  
plt.ylim(x2.min(), x2.max())  
  
for i, j in enumerate(np.unique(y_set)):
```

```
plt.scatter(  
    x_set[y_set == j, 0],  
    x_set[y_set == j, 1],  
    c=ListedColormap(('purple', 'green'))(i),  
    label=j  
)
```

```
plt.title('Decision Tree (Training Set)')
```

```
plt.xlabel('Age')
```

```
plt.ylabel('Estimated Salary')
```

```
plt.legend()
```

```
plt.show()
```



## Testing set visualization:

```
from matplotlib.colors import ListedColormap
```

```
x_set, y_set = x_test, y_test
x1, x2 = np.meshgrid(
    np.arange(x_set[:, 0].min() - 1, x_set[:, 0].max() + 1, 0.01),
    np.arange(x_set[:, 1].min() - 1, x_set[:, 1].max() + 1, 0.01)
)

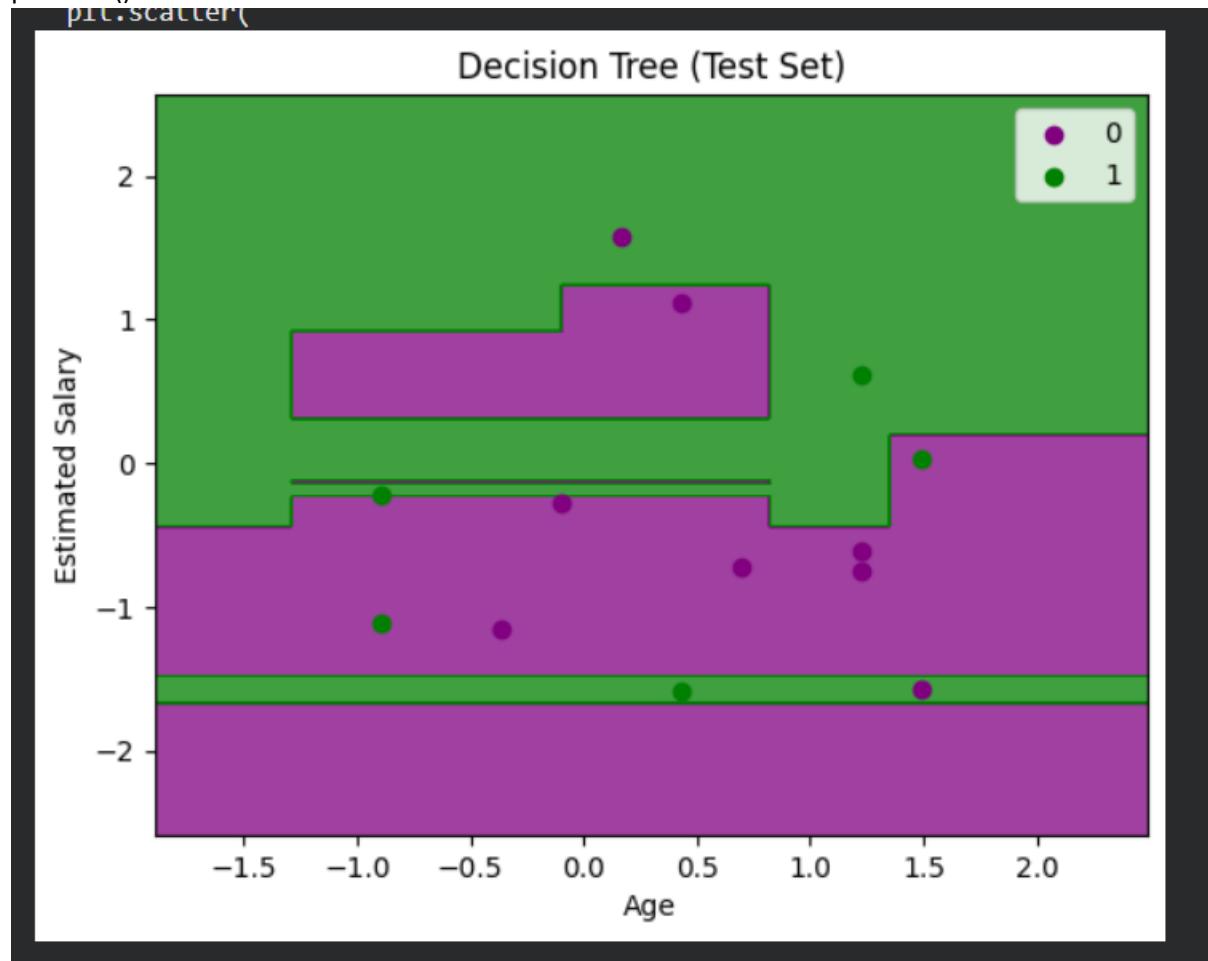
plt.contourf(
    x1, x2,
    classifier.predict(np.array([x1.ravel(), x2.ravel()]).T).reshape(x1.shape),
    alpha=0.75,
    cmap=ListedColormap(('purple', 'green'))
)

plt.xlim(x1.min(), x1.max())
plt.ylim(x2.min(), x2.max())

for i, j in enumerate(np.unique(y_set)):
    plt.scatter(
        x_set[y_set == j, 0],
        x_set[y_set == j, 1],
        c=ListedColormap(('purple', 'green'))(i),
        label=j
    )

plt.title('Decision Tree (Test Set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
```

```
plt.show()  
plt.scatter(
```



# Experiment 10:

## Creating dataset:

```
import pandas as pd  
  
import numpy as np  
  
  
data = {  
  
    "feat1": [5,6,4,5,7,8,5,6,4,5,7,8, 10,11,9,12,13,11,10,12,9,13,12,11,  
              20,21,19,18,22,23,19,20,21,22,18,23, 6,7,5,8,9,7,6,8,5,9,7,8,  
              14,15,13,16,17,15,14,16,13,17,15,16],  
  
  
    "feat2": [10,11,12,9,8,12,11,10,9,8,13,12, 20,22,19,21,23,20,22,21,19,23,22,21,  
              30,29,28,31,32,30,29,31,28,32,30,29, 15,16,14,17,18,16,15,17,14,18,16,17,  
              25,26,24,27,28,26,25,27,24,28,26,27],  
  
  
    "feat3": [2,3,1,2,4,5,2,3,1,2,4,5, 6,7,5,8,9,7,6,8,5,9,7,8,  
              1,2,3,1,4,5,3,2,1,4,3,5, 7,8,6,9,10,8,7,9,6,10,8,9,  
              3,4,2,5,6,4,3,5,2,6,4,5],  
  
  
    "feat4": [1,1,2,1,3,2,1,2,1,3,2,1, 4,4,5,4,6,5,4,5,4,6,5,4,  
              7,8,6,7,9,8,7,8,6,9,8,7, 2,3,1,2,4,3,2,3,1,4,3,2,  
              5,6,4,5,7,6,5,6,4,7,6,5],  
  
}  
  
  
labels = [0]*12 + [1]*12 + [2]*12 + [0]*12 + [1]*12 # 60 labels total  
  
  
df = pd.DataFrame(data)  
df["label"] = labels  
  
  
X = df[["feat1", "feat2", "feat3", "feat4"]]  
y = df["label"]
```

```
df.head()
```

### KNN Classification:

```
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.model_selection import train_test_split  
import matplotlib.pyplot as plt
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

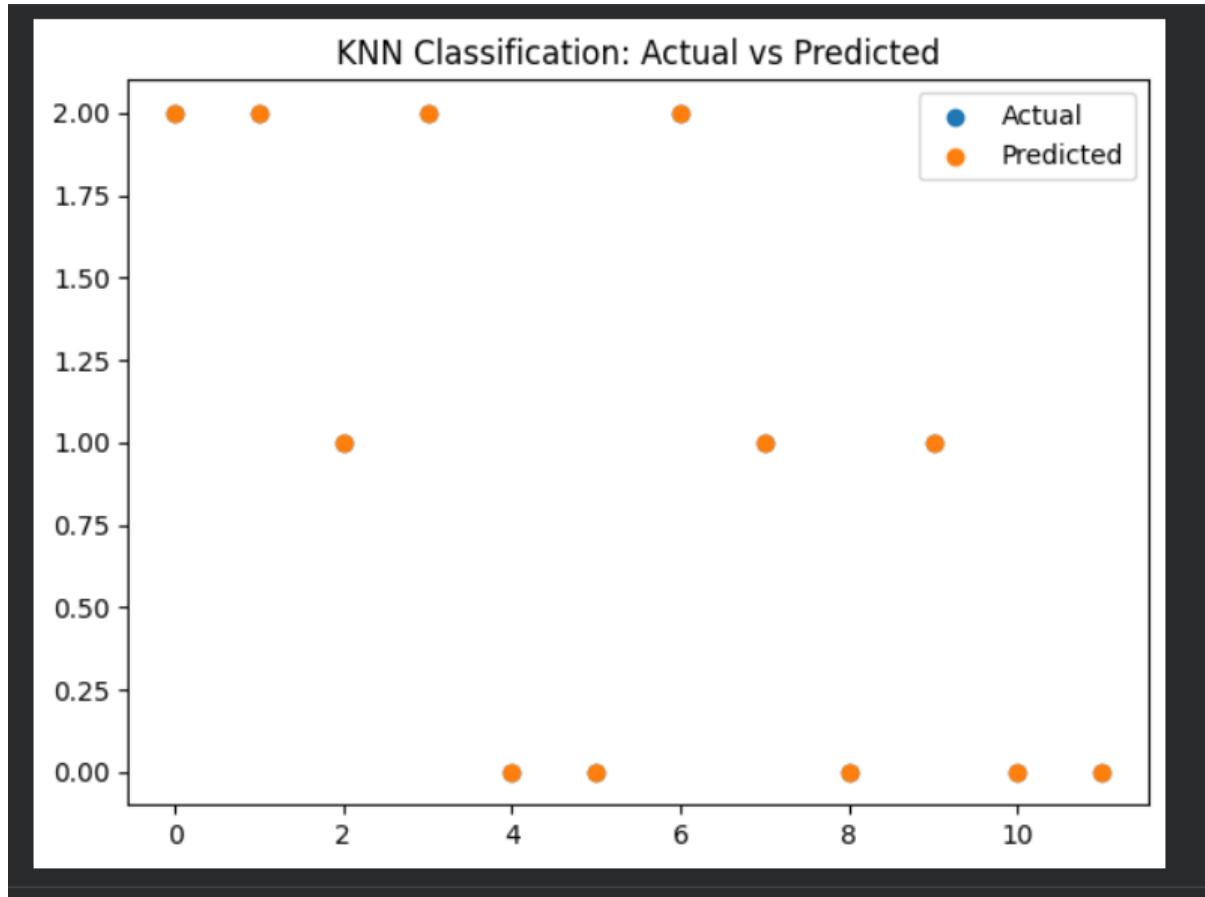
```
knn = KNeighborsClassifier(n_neighbors=3)
```

```
knn.fit(X_train, y_train)
```

```
y_pred = knn.predict(X_test)
```

```
plt.figure(figsize=(7,5))  
plt.scatter(range(len(y_test)), y_test, label="Actual")  
plt.scatter(range(len(y_pred)), y_pred, label="Predicted")  
plt.title("KNN Classification: Actual vs Predicted")  
plt.legend()
```

```
plt.show()
```



# Experiment 11:

## Creating dataset:

```
import pandas as pd
```

```
import numpy as np
```

```
data = {
```

```
    "feat1": [5,6,4,5,7,8,5,6,4,5,7,8, 10,11,9,12,13,11,10,12,9,13,12,11,  
              20,21,19,18,22,23,19,20,21,22,18,23, 6,7,5,8,9,7,6,8,5,9,7,8,  
              14,15,13,16,17,15,14,16,13,17,15,16],
```

```
    "feat2": [10,11,12,9,8,12,11,10,9,8,13,12, 20,22,19,21,23,20,22,21,19,23,22,21,  
              30,29,28,31,32,30,29,31,28,32,30,29, 15,16,14,17,18,16,15,17,14,18,16,17,  
              25,26,24,27,28,26,25,27,24,28,26,27],
```

```
    "feat3": [2,3,1,2,4,5,2,3,1,2,4,5, 6,7,5,8,9,7,6,8,5,9,7,8,  
              1,2,3,1,4,5,3,2,1,4,3,5, 7,8,6,9,10,8,7,9,6,10,8,9,  
              3,4,2,5,6,4,3,5,2,6,4,5],
```

```
    "feat4": [1,1,2,1,3,2,1,2,1,3,2,1, 4,4,5,4,6,5,4,5,4,6,5,4,  
              7,8,6,7,9,8,7,8,6,9,8,7, 2,3,1,2,4,3,2,3,1,4,3,2,  
              5,6,4,5,7,6,5,6,4,7,6,5],
```

```
}
```

```
labels = [0]*12 + [1]*12 + [2]*12 + [0]*12 + [1]*12 # 60 labels total
```

```
df = pd.DataFrame(data)
```

```
df["label"] = labels
```

```
X = df[["feat1", "feat2", "feat3", "feat4"]]
```

```
y = df["label"]
```

```
df.head()
```

## K - Means Clustering:

```
from sklearn.cluster import KMeans
```

```
kmeans = KMeans(n_clusters=3, random_state=0)
```

```
clusters_km = kmeans.fit_predict(X)
```

```
plt.figure(figsize=(7,5))
```

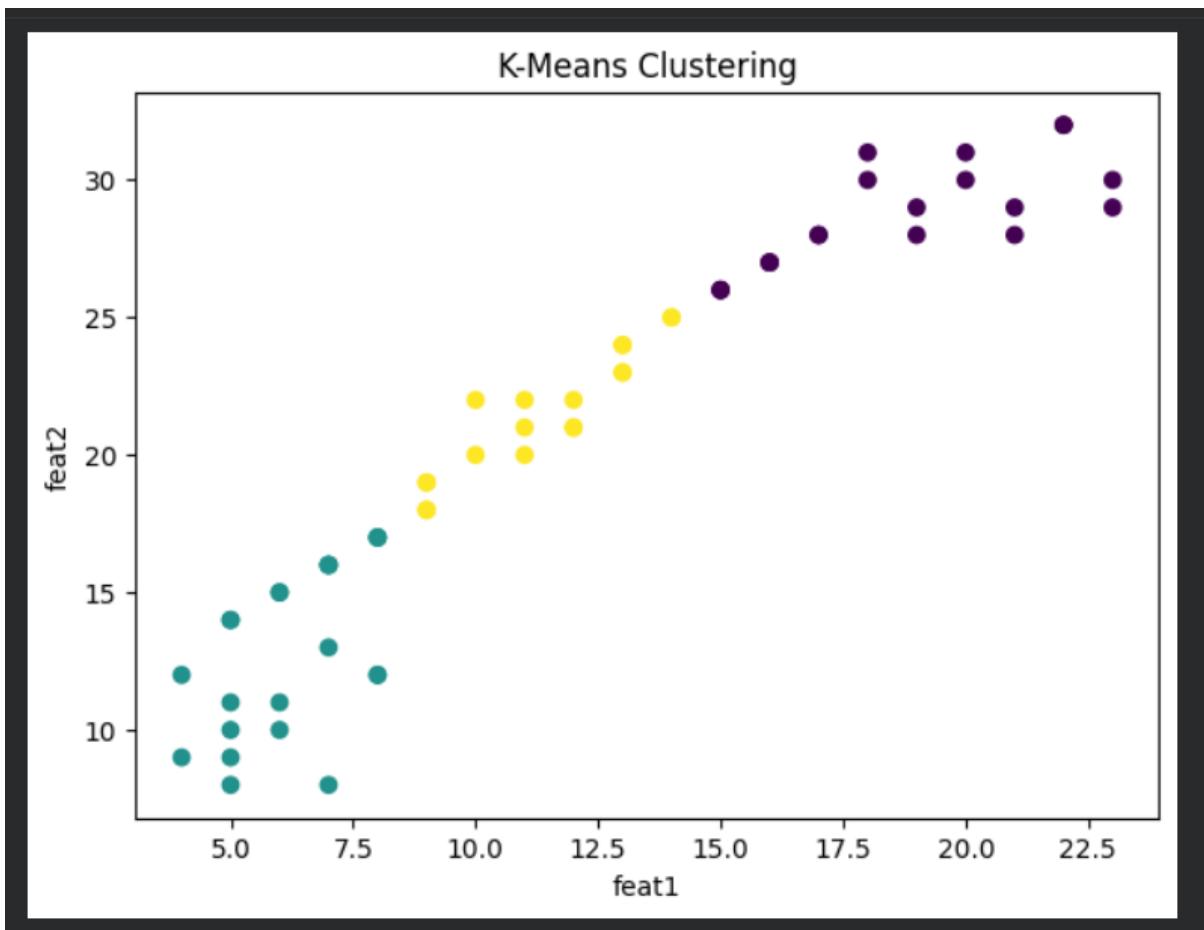
```
plt.scatter(X["feat1"], X["feat2"], c=clusters_km)
```

```
plt.title("K-Means Clustering")
```

```
plt.xlabel("feat1")
```

```
plt.ylabel("feat2")
```

```
plt.show()
```



# Experiment 12:

## Creating dataset:

```
import pandas as pd  
  
import numpy as np  
  
  
data = {  
  
    "feat1": [5,6,4,5,7,8,5,6,4,5,7,8, 10,11,9,12,13,11,10,12,9,13,12,11,  
              20,21,19,18,22,23,19,20,21,22,18,23, 6,7,5,8,9,7,6,8,5,9,7,8,  
              14,15,13,16,17,15,14,16,13,17,15,16],  
  
    "feat2": [10,11,12,9,8,12,11,10,9,8,13,12, 20,22,19,21,23,20,22,21,19,23,22,21,  
              30,29,28,31,32,30,29,31,28,32,30,29, 15,16,14,17,18,16,15,17,14,18,16,17,  
              25,26,24,27,28,26,25,27,24,28,26,27],  
  
    "feat3": [2,3,1,2,4,5,2,3,1,2,4,5, 6,7,5,8,9,7,6,8,5,9,7,8,  
              1,2,3,1,4,5,3,2,1,4,3,5, 7,8,6,9,10,8,7,9,6,10,8,9,  
              3,4,2,5,6,4,3,5,2,6,4,5],  
  
    "feat4": [1,1,2,1,3,2,1,2,1,3,2,1, 4,4,5,4,6,5,4,5,4,6,5,4,  
              7,8,6,7,9,8,7,8,6,9,8,7, 2,3,1,2,4,3,2,3,1,4,3,2,  
              5,6,4,5,7,6,5,6,4,7,6,5],  
}  
  
labels = [0]*12 + [1]*12 + [2]*12 + [0]*12 + [1]*12 # 60 labels total  
  
df = pd.DataFrame(data)  
df["label"] = labels  
  
X = df[["feat1", "feat2", "feat3", "feat4"]]
```

```
y = df["label"]
```

```
df.head()
```

## DBSCAN Clustering:

```
from sklearn.cluster import DBSCAN
```

```
db = DBSCAN(eps=3, min_samples=3)
```

```
clusters_db = db.fit_predict(X)
```

```
plt.figure(figsize=(7,5))
```

```
plt.scatter(X["feat1"], X["feat2"], c=clusters_db)
```

```
plt.title("DBSCAN Clustering")
```

```
plt.xlabel("feat1")
```

```
plt.ylabel("feat2")
```

```
plt.show()
```

DBSCAN Clustering

